

Chi Squared Tests: Goodness of fit, Independence


9.1, 9.2

Today's topics

Chi Squared Tests

- Goodness of fit (9.1)
- Independence (9.2)
- Exact p-value

Response / Dependent
categorical data



χ^2 Goodness of Fit Test - Background

/ categorical var

Developed by Pearson in 1900

Tests the appropriateness of different models

Approximate test for use with large samples (large n).

Only appropriate when all expected cell counts are greater than 5. (Otherwise, should consider calculating exact p -value)



χ^2 Goodness of Fit Test

20 C ← 30 D ←

10 other ←

group

$k=3$

Random sample of size n is classified into k categories or cells.

Let Y_1, Y_2, \dots, Y_k denote the respective cell frequencies;

$$\sum_{i=1}^k Y_i = n$$

Denote cell probabilities p_1, p_2, \dots, p_k .

→ $H_0: p_1 = p_{10}, p_2 = p_{20}, \dots, p_k = p_{k0}$

p_{i0}

$H_A: H_0$ not true

p_{20}

numerical

1-prop z test
 $H_0: p = .2$
 $H_A: p \neq .2$

$$np \quad \left\{ \quad np_1, np_2, np_3, \dots \right.$$

χ^2 Goodness of Fit Test

	Group 1	Group 2	...	Group k	Total
Observed Freq (O_i)	Y_1	Y_2	...	Y_k	n
Probability H_0 (p_{i0})	p_{10}	p_{20}	...	p_{k0}	1
Expected Freq (E_i)	np_{10}	np_{20}	...	np_{k0}	n

always right tailed

C	D	O
~	~	~

χ^2 Goodness of Fit Test

$$df = 3 - 1 = 2$$

→ Test statistic:

$$X^2 = \sum_{i=1}^k \frac{(Y_i - np_{i0})^2}{np_{i0}} = \underbrace{\sum_{i=1}^k \frac{(Obs_i - Exp_i)^2}{Exp_i}}_{\text{test stat}} \sim \chi^2_{(k-1)}$$

Data ↓

Reject H_0 if $X^2 \geq \underbrace{\chi^2_{(k-1), \alpha}}_{\text{crit}}$



$$\sum \frac{(O - \bar{E})^2}{\bar{E}}$$

Random Digit Example (Goodness of Fit)

When making random numbers, people are usually reluctant to record the same or consecutive numbers in adjacent positions. Even though these true probabilities should be $p_{10} = 1/10$ and $p_{20} = 2/10$, respectively. gt8suv tests a friend's concept of a random sequence by asking them to generate a sequence of 51 *random* digits.

May0 generates the following sequence:

(Ex 9.1-1)

5	8	3	1	9	4	6	7	9	2	6	3	0
8	7	5	1	3	6	2	1	9	5	4	8	0
3	7	1	4	6	0	4	3	8	2	7	3	9
8	5	6	1	8	7	0	3	5	2	5	2	

Use a goodness of fit test to make a statistical decision about whether this sequence seems to be truly random or not (at $\alpha = 0.05$)

50

5	8	3	1	9	4	6	7	9	2	6	3	0
8	7	5	1	3	6	2	1	9	5	4	8	0
3	7	1	4	6	0	4	3	8	2	7	3	9
8	5	6	1	8	7	0	3	5	2	5	2	

Random Digit Example (Goodness of Fit)

$$H_0: p_1 = p_{10} = 1/10, p_2 = p_{20} = 2/10, p_3 = p_{30} = 7/10$$

$$H_A: ?$$

Group 1 G2 G3
Same cons other

	Group 1	Group 2	...	Group k	Total
Observed Freq (O_i)	Y_1 0	Y_2 8	...	Y_k 42	n 50
Probability H_0 (p_{i0})	p_{10} 1/10	p_{20} 2/10	...	p_{k0} 7/10	1
Expected Freq (E_i)	np_{10} 5	np_{20} 10	...	np_{k0} 35	n

test stat ✓

$$\frac{(0-5)^2}{5} + \frac{(8-10)^2}{10} + \frac{(42-35)^2}{35} = 6.8$$

$$6.8 > 5.991 = \chi_{0.05}^2(2).$$



$$\sum \frac{(O-E)^2}{E} = 0$$

Small



- Reject H_0 :
- There is significant evidence at $\alpha = 0.05$ to suggest a lack of fit. The sequence does **not** seem to be random.

* good fit

Obs

~~8~~ 4

~~10~~ 11

35

Exp

5

10

35

3 factors
or levels

P_{10}

	Y	B	R
O	25	40	35
E	30	50	20

Quick Example Goodness of fit

Claim

- ~~M&Ms~~ ^{Skittles}: 1, 2, 3
- 30% yellow, 50% blue, 20% red

- You open a bag of 100 ~~M&Ms~~ ^{Skittles} and get 25 Y, 40 B, 35 R

$$H_0: P_1 = .3, P_2 = .5, P_3 = .2$$

1 variable: color

$$\sum \frac{(O-E)^2}{E} = \frac{5^2}{30} + \frac{10^2}{50} + \frac{15^2}{20} \sim \chi^2_2$$

Definition

→ Contingency table

Summarizes relationship between categorical variables by displaying frequency distribution.

e.g.

Table 9.2-1 Undergraduates at the University of Iowa						
Gender	College					Totals
	Business	Engineering	Liberal Arts	Nursing	Pharmacy	
Male	21	16	145	2	6	190
Female	14	4	175	13	4	210
Totals	35	20	320	15	10	400

χ^2 Test for Homogeneity & Independence (9.2)

same test

Tests relationship between two categorical variables.

The difference in the two names depends on how the data is collected.

χ^2 Test for Homogeneity & Independence (9.2)

χ^2 Test for **Homogeneity**: (one margin fixed)

Tests whether two or more sub-groups of a population share the same distribution of a single categorical variable. E.g., do different age groups have the same proportion of people who prefer Twitch, YouTube Live, or Zoom?

Independent random samples from r populations.

Each sample is classified into c response categories.

H_0 : In each category, the probabilities are equal for all r populations.

A hand-drawn contingency table on a dark background. The columns are labeled 'AGE' with categories 'Y' (Young), 'M' (Middle), and 'O' (Old). The rows are labeled 'Pre' (Preference) with categories 'T' (Twitch), 'Y' (YouTube Live), and 'Z' (Zoom). The table is a 3x3 grid with yellow lines. There are also some additional handwritten marks: a '5' next to the 'Pre' label, and a '13' in the bottom right corner of the grid.

	Y	M	O
Pre 5			
T			
Y			
Z			

χ^2 Test for Homogeneity & Independence (9.2)

 χ^2 Test for **Independence**: (no margins fixed)

data collection step

Tests whether two categorical variables are associated with one another in the population, e.g. age group vs video streaming platform preference.

A random sample of size n is simultaneously classified with respect to two characteristics, one has r categories and the other c categories.

H_0 : The two classifications are independent; i.e., each cell probability is the product of the row and column marginal probabilities

(9.2) notes

2 cat.
variables

Grade, Instructor

Homogeneity (Independence)

In each group (BooleanHypercube, Obese future), I collect a separate sample. I look at whether Instructor and grade are related

Independence

I collect a large sample of students who took 420 and ask them who they had as an instructor, as well as their grade.

notes

Why would we do a test for homogeneity?

Say you have a rare disease that you are trying to determine something about

If you do test for independence and collect a sample of size n

χ^2 Test for Homogeneity & Independence (9.2)

Test statistic (both tests):

$$\chi^2 = \sum_{\text{cells}} \frac{(O-E)^2}{E} \sim \chi^2_{(r-1)*(c-1)}$$

O = observed cell frequency

Handwritten example of a 2x2 contingency table with marginal totals:

1	9	10
9	1	10
10	10	20

Handwritten notes: $(2-1) \cdot (2-1) = 1$ df, $10 \cdot 10 / 20 = 5$

$$E = \frac{\text{row total} * \text{column total}}{\text{overall total}}$$

Handwritten example of a 3x3 contingency table with symbols:

#	#	F
F	F	
#	#	

(more formally):

$$\chi^2 = \sum_{j=1}^h \sum_{i=1}^k \frac{(y_{ij} - n_{.j} p_{i.})^2}{n_{.j} p_{i.}} \sim \chi^2_{(r-1)(c-1)}$$

χ^2 Test of homogeneity example:

Saumdog is wondering which instructor to select for Stat 420. Albert doesn't know much so he states H_0 : their grade distributions are the same.

We collect 2 separate random samples from each instructor and obtain the following data.

Instructor	Grade					Totals
	A	B	C	D	F	
ObeseFuture	8	13	16	10	3	50
Boolean HyperCube	4	9	14	16	7	50

Instructor

Perform a χ^2 test of homogeneity at significance level 0.05 to determine if these grade distributions are similar.

$$\alpha = 0.05$$

obs

Instructor	Grade					Totals
	A	B	C	D	F	
ObeseFuture	8	13	16	10	3	50
Boolean HyperCube	4	9	14	16	7	50

$$\frac{50 \times 12}{100} = 6$$

$$\frac{30 \times 50}{100} = 15$$

EXP

$$(5-1)(2-1) = 4$$

$$12 \quad 22 \quad 30 \quad 26 \quad 10 \quad 100$$

$\chi^2 =$

$$= \frac{(8-6)^2}{6} + \frac{(13-11)^2}{11} + \frac{(16-15)^2}{15} + \frac{(10-13)^2}{13} + \frac{(3-5)^2}{5} + \frac{(4-6)^2}{6} + \frac{(9-11)^2}{11} + \frac{(14-15)^2}{15} + \frac{(16-13)^2}{13} + \frac{(7-5)^2}{5}$$

$$= \frac{4}{6} + \frac{4}{11} + \frac{1}{15} + \frac{9}{13} + \frac{4}{5} + \frac{4}{6} + \frac{4}{11} + \frac{1}{15} + \frac{9}{13} + \frac{4}{5} = 5.18$$

test stat

χ^2_4

• Test of independence

distrib



say

I

collect


$n = 100$

students

Homogeneity ←

	Grade					Totals
	A	B	C	D	F	
ObeseFuture	8	13	16	10	3	50
Boolean HyperCube	4	9	14	16	7	50

Inst



	Grade					Tot
	A	B	C	D	F	
OF	12	10	14	10	10	60
BH	8	6	10	8	8	40
	20	16	24	18	18	100

notes

DN12 H₀

5.18

χ^2_4






Table IV The Chi-Square Distribution



$$P(X \leq x) = \int_0^x \frac{1}{\Gamma(r/2)2^{r/2}} w^{r/2-1} e^{-w/2} dw$$

	$P(X \leq x)$							
	0.010	0.025	0.050	0.100	0.900	0.950	0.975	0.990
r	$\chi^2_{0.99}(r)$	$\chi^2_{0.975}(r)$	$\chi^2_{0.95}(r)$	$\chi^2_{0.90}(r)$	$\chi^2_{0.10}(r)$	$\chi^2_{0.05}(r)$	$\chi^2_{0.025}(r)$	$\chi^2_{0.01}(r)$
1	0.000	0.001	0.004	0.016	2.706	3.841	5.024	6.635
2	0.020	0.051	0.103	0.211	4.605	5.991	7.378	9.210
3	0.115	0.216	0.352	0.584	6.251	7.815	9.348	11.34
4	0.297	0.484	0.711	1.064	7.779	9.488	11.14	13.28

5.18
t.s.
9.488
C.V.
↑

Example


A random sample of 400 undergraduate students at the University of Iowa was selected, then classified by college and gender. Are these variables independent at $\alpha = 0.01$?

Table 9.2-1 Undergraduates at the University of Iowa

Gender	College					Totals
	Business	Engineering	Liberal Arts	Nursing	Pharmacy	
Male	21	16	145	2	6	190
Female	14	4	175	13	4	210

Table 9.2-1 Undergraduates at the University of Iowa

Gender	College					Totals
	Business	Engineering	Liberal Arts	Nursing	Pharmacy	
Male	21 (16.625)	16 (9.5)	145 (152)	2 (7.125)	6 (4.75)	190
Female	14 (18.375)	4 (10.5)	175 (168)	13 (7.875)	4 (5.25)	210
Totals	35	20	320	15	10	400



$$\frac{(21 - 16.625)^2}{16.625} + \frac{(14 - 18.375)^2}{18.375} + \dots + \frac{(4 - 5.25)^2}{5.25}$$
$$1.15 + 1.04 + 4.45 + 4.02 + 0.32 + 0.29 + 3.69$$
$$+ 3.34 + 0.33 + 0.30 = 18.93.$$

critical
value : 13.28

Exact p-value

(not really a new definition) ←

Review: What is the definition of a p-value?

Calculate this probability directly, instead of using a Normal or Chi Squared approximation.

prop test

$$H_0: p = 0.5$$

$$H_A: p > 0.5$$

Type II Error

Exact p-value - Example

Luigi wants to test whether a coin is a fair coin or not (loaded in favor of Heads) at $\alpha = 0.05$. Data from 100 flips (heads = 1, tails = 0):

```
a = rbinom(100, 1, 0.6)
mean(a) #0.58
```

1 1 1 0 1 0 0 0 1 0 1 1 1 0 1 0 0 1 0 0 0 0 1 0 1 0 1 1 1 0 0 1 0 1 1 1 1 0 0 1 1 1 1 0 0 1 0 0 1 0 1 1 1 1 0 0 1 1 1 1 0 0 1 1 1 1 0 0 1 1 1 1 0 0 1 1 1 1 0 0 1 1 1 1 1 0 1 1 1 1 1 0 1 1 1 1 0 1 0 1 0 1

(Could find Z value and do a test that we already know)

OR (9.1) \rightarrow 58 1's
42 0's

To find exact p-value, just use the definition of the p-value

Let X = # of heads in 100 tosses Under $H_0: X \sim \text{Binomial}(100, 0.5)$

P-value = $P[X \geq 58 | p = 0.5]$.

$$= 1 - P[X \leq 57 | p = 0.5]$$

In R:

```
> 1 - pbinom(57, 100, .5)
[1] 0.06660531
```

exact p-value

notes

if $n = 10000$
approx easier

if small sample sizes,

can't use CLT, can't use χ^2

use exact p-value

notes

if α was lower 0.01
what would happen to Type II error?



$\alpha \uparrow$, $\beta \downarrow$